Data Science TA Hiring Assignment

Bonus Section

Rohan Saha

April 10, 2025

Supporting Students with Struggling Concepts or Deadlines

As a Teaching Assistant for Data Science, my approach to supporting students who are struggling with concepts or deadlines would be multi-faceted and personalized:

For Conceptual Struggles

I would first identify the specific knowledge gaps through one-on-one discussions. Understanding whether the challenge is with foundational knowledge, application, or advanced concepts allows me to tailor my approach.

For visual learners, I would prepare simplified diagrams and flowcharts that break down complex ideas into digestible components. For example, explaining Random Forest by first showing a simple decision tree, then demonstrating how multiple trees work together.

I believe in the power of real-world examples. I would connect abstract statistical concepts to tangible scenarios relevant to students' interests. For instance, explaining regression by analyzing factors affecting movie success if a student is interested in film.

Supplementary resources would be curated based on different learning styles—video tutorials for visual learners, interactive notebooks for hands-on learners, and concise written summaries for those who prefer reading.

For Deadline Challenges

I would help students develop a structured plan by breaking large assignments into smaller, manageable milestones with self-imposed deadlines before the final submission date.

Regular check-ins would be established to monitor progress and provide early intervention if a student is falling behind.

I recognize that mental blocks often cause delays. I would teach time management techniques and encourage brief breaks to refresh their perspective when stuck.

Most importantly, I would create a supportive environment where students feel comfortable discussing challenges without judgment, emphasizing that seeking help is a strength,

not a weakness.

Breaking Down "Gradient Descent" for Beginners

Step 1: The Mountain Analogy

Imagine you're standing on a mountain in thick fog, and your goal is to reach the lowest point in the valley. Since you can't see the whole landscape, how would you proceed? You would feel the ground around you and take a step in the direction where the slope goes downward. After taking that step, you'd repeat the process—feel the ground and step downward again.

This is exactly what gradient descent does! It's an algorithm that helps us find the minimum value of a function when we can't see the entire function at once.

Step 2: Connecting to Machine Learning

In machine learning, our "mountain" is the error function—how wrong our model's predictions are. Our goal is to find the model parameters (like weights in a neural network) that make this error as small as possible.

Step 3: The Mathematics, Simplified

Let's break it down with a simple example:

- We have a model with one parameter (w)
- We calculate the error for our current parameter value
- We find the derivative (slope) of the error function at our current position
- We update our parameter: new w = old w - (learning rate × slope)
- The learning rate determines how big our steps are—too large and we might overshoot; too small and we'll take forever

Step 4: Visual Understanding

I would draw a simple bowl-shaped function on the board and trace the path of gradient descent, showing how it zigzags toward the bottom with each step, gradually getting closer to the minimum.

Step 5: Hands-on Mini-Example

To solidify understanding, I would walk through a tiny example with actual numbers where we start with w=2, calculate the error and its derivative, update w with a learning rate of 0.1, and observe how the error decreases.

The key is to present gradient descent not as an intimidating mathematical concept, but as a logical approach to solving a problem—similar to how we naturally find our way down a mountain in the fog.